

! This class has been made inactive. No posts will be allowed until an instructor reactivates the class.

note

172 views

Sharing "StrategyLearner" from mc3-p3

Professor mentioned that is now ok to share all the way to MC3-P3, if you would like to share your solution for MC3-P3 please post here.

NB: Professor mentioned that he was ok with code being shared here (as copy&paste here and not to link to it from an external website).

Thanks in advance for anyone sharing!

mc3-p3

Updated 1 year ago by Carlos Aguayo

followup discussions for lingering questions and comments

☒ Resolved ☐ Unresolved



Carlos Aguayo 1 year ago

Here's my solution:

```
"""
Template for implementing StrategyLearner (c) 2016 Tucker Balch

StrategyLearner
Carlos Aguayo
carlos.aguayo@gatech.edu
gtid 903055858

"""

import datetime as dt
import QLearner as ql
import pandas as pd
import util as ut
import numpy as np

import warnings # http://stackoverflow.com/questions/15777951/how-to-suppress-pandas-future-warning
warnings.simplefilter(action="ignore", category=FutureWarning)

class StrategyLearner(object):

    # constructor
    def __init__(self, verbose = False):
        self.verbose = verbose
        self.learner = None

    @staticmethod
    def bollinger_value(price, sma, stddev):
        # bb_value[t] = (price[t] - SMA[t]) / (2 * stddev[t])
        return (price - sma) / (2 * stddev)

    @staticmethod
    def momentum(price, n):
        # momentum[t] = (price[t]/price[t-N]) - 1
        return price / price.shift(n) - 1

    @staticmethod
    def volatility(price, n):
        # Volatility is just the stdev of daily returns.

        def compute_daily_returns(s):
            # https://www.udacity.com/course/viewer#!/c-ud501/L-4156938722/e-4185858982/m-4185858984
            daily_returns = s.copy()
            daily_returns[1:] = (s[1:] * 1.0 / s[:-1].values) - 1
            return daily_returns[1:]

        return pd.rolling_std(compute_daily_returns(price), n)

    @staticmethod
    def get_bins(data, steps):
        data = data.copy()
        steps -= 1

        stepsize = len(data) / steps
        data.sort()
        threshold = [0] * steps
        for i in range(steps):
            threshold[i] = data[(i + 1) * stepsize]
        return threshold
```

```

@staticmethod
def to_technical_features(prices, symbol):
    # Bollinger Bands, Momentum, Volatility
    time_window = 20

    sma = pd.rolling_mean(prices, time_window)
    stddev = pd.rolling_std(prices, time_window)

    data = np.ndarray([len(prices), 3])

    data[:, 0] = StrategyLearner.bollinger_value(price=prices, sma=sma, stddev=stddev)[symbol]
    data[:, 1] = StrategyLearner.momentum(price=prices, n=5)[symbol]
    data[:, 2] = StrategyLearner.volatility(price=prices, n=time_window)[symbol]
    data[:, :] = np.nan_to_num(data)

    return data

def build_state(self, array_technical_features, holding):
    state = "".join(map(lambda discrete_value: str(int(discrete_value)), array_technical_features))
    return int(state + str(holding))

# this method should create a QLearner, and train it for trading
def addEvidence(self,
                symbol="IBM",
                sd=dt.datetime(2008, 1, 1),
                ed=dt.datetime(2009, 1, 1),
                sv=10000):

    num_states = 10**4 # 3 features with 9 values, times 3 because buy/sell/hold

    # https://piazza.com/class/ij9yiif53l27fs?cid=1640
    self.learner = ql.QLearner(num_states=num_states,
                                num_actions=3, # 3 actions, buy, sell or hold
                                alpha=0.2,
                                gamma=0.9,
                                rar=0.98,
                                radr=0.999,
                                dyna=0,
                                verbose=False)

    # example usage of the old backward compatible util function
    syms=[symbol]
    dates = pd.date_range(sd, ed)
    df = ut.get_data(syms, dates) # automatically adds SPY
    prices = df[syms] # only portfolio symbols

    # example use with new colname
    # volume_all = ut.get_data(syms, dates, colname="Volume") # automatically adds SPY
    # volume = volume_all[syms] # only portfolio symbols
    # volume_SPY = volume_all['SPY'] # only SPY, for comparison later
    # if self.verbose: print volume

    data = self.to_technical_features(prices, symbol)

    # TODO - Can't I do in one line?
    # TODO - I'd need to store the bins right?
    for i in range(data.shape[1]):
        data[:, i] = np.digitize(x=data[:, i], bins=self.get_bins(data[:, i], 10))

    round_lot = 100
    max_iterations = 100

    cumulative_return = np.ndarray(max_iterations)

    for iteration in range(0, max_iterations):
        # holding - 0
        # buy - 1
        # sell - 2
        df['cash'] = 0
        df['portfolio_value'] = 0
        df['cash'].ix[0] = sv
        df['portfolio_value'].ix[0] = sv

        initial_state = self.build_state(data[0], 0)
        action = self.learner.querysetstate(initial_state)
        prev_date = df.index[0]

        i = 1
        position = 0
        for date in df.index[1:]:
            entry = data[i]
            i += 1

            today_price = prices.ix[date, symbol]

            invalid = False

            if action == 1: # buy
                position += 1
                if position > 1:

```

```

        position = 1
        invalid = True
    elif action == 2: # sell
        position -= 1
        if position < -1:
            position = -1
            invalid = True

    if invalid:
        action = 0 # TODO: Can I use something like an enum instead of 0, 1 and 2?

    if action == 0: # hold
        df.ix[date, 'cash'] = df.ix[prev_date, 'cash']
    elif action == 1: # buy
        df.ix[date, 'cash'] = df.ix[prev_date, 'cash'] - today_price * round_lot

```



Alex K 1 year ago Thanks for sharing.

☒ Resolved ☐ Unresolved



Andrew 11 months ago

I was hoping to get confirmation on the discrepancy in the feedback before posting this. However, with the class over, I get the sense that everyone is moving on, which is understandable. I have resolved what I believe was the source to my own satisfaction, so I'll go ahead and post it anyways, as a few have asked for it.

A few caveats, of course. As I've mentioned elsewhere, I'm not a "Python guy", so I'm sure there are better ways of doing some of this. I wouldn't even claim that this is the best way to do the project. It's also more than a bit messy - pieces of this grew organically over the last few projects and I didn't bother to clean things up much once I got it working well enough. The Indicators class in particular needs a rewrite, but you can at least see a number of the features I tried at one time or another on past projects. Of course, if I had to write it all over again I'd probably do everything differently anyways.

I've gone back and sprinkled some comments throughout. If anything isn't clear, feel free to ask questions.

```

"""
    StrategyLearner: Uses Q-Learning to learning trading policy.
"""

import datetime as dt
import math

import numpy as np
import pandas as pd

import QLearner as ql
from util import get_data

class StrategyLearner(object):

    def __init__(self, verbose=False):
        self._verbose = verbose
        self._q_learner = None

        # indicator features to use
        self._features = ["momentum", "momentum_10", "volume", "volatility", "10_day", "20_day", "5_day_spy"]
        # extra fields to pull from data (besides the defaults)
        self._fields = ['High', 'Low', 'Close', 'Volume']
        # support transactions with 100 or 200 shares
        self._shares = [100, 100, 200, 200, None]
        # actions: buy/sell 100 shares, buy/sell 200 shares, hold
        self._actions = ["BUY", "SELL", "BUY", "SELL", "HOLD"]

        self._indicators = None
        self._portfolio = None

    def addEvidence(self, symbol="IBM", sd=dt.datetime(2007, 12, 31), ed=dt.datetime(2009, 12, 31), sv=10000):

        # get data for our states (training period):
        # states - chronological list of indicator state-strings
        # state_map - map of state-strings to states
        # dates - dates for these states
        # price_dict - full price data as dictionary
        (states, state_map, dates, price_dict) = self._get_states_data(symbol, sd, ed)

        # count the total number of states
        num_states = len(state_map.keys())

        self._print_debug("Number of reachable states (train): " + str(num_states))

        # set up the Q-Learner
        self._q_learner = ql.QLearner(num_states, len(self._actions), rar=0.98, radr=0.999)

        # variables to track stopping criteria
        cr_last = None
        last_change = 0
        iterations = 0

```

```

# continue to perform Q-Learning iterations until stopping criteria for convergence
while True:
    iterations += 1
    # perform trades and get the current cumulative returns
    cr = self._perform_trades(price_dict, symbol, sd, sv, dates, states, state_map)

    self._print_debug("Cumulative return at iteration " + str(iterations) + ": " + str(cr))

    # if the CR has changed, update the last stored value
    if not cr_last or (math.fabs(cr_last - cr) >= 0.001):
        cr_last = cr
        last_change = iterations

    # check the stopping criteria (CR hasn't changed for at least 10 iterations)
    if iterations - last_change > 10:
        break

self._print_debug("Final training results: " + str(cr_last) + " after " + str(iterations) + " iterations")

def testPolicy(self, symbol="IBM", sd=dt.datetime(2009, 12, 31), ed=dt.datetime(2011, 12, 31), sv=10000):

    # get data for our states (test period):
    # states - chronological list of indicator state-strings
    # state_map - map of state-strings to states
    # dates - dates for these states
    # price_dict - full price data as dictionary
    (states, state_map, dates, price_dict) = self._get_states_data(symbol, sd, ed, self._indicators)

    # count the total number of states
    num_states = len(state_map.keys())

    self._print_debug("Number of reachable states (test): " + str(num_states))

    # perform test trades for this period
    cr = self._perform_trades(price_dict, symbol, sd, sv, dates, states, state_map, True)

    self._print_debug("Cumulative return: " + str(cr))

    # return the generated trades for this period
    return self._portfolio.trades

def _get_states_data(self, symbol, sd, ed, indicators=None):
    # get data for symbol
    dates = pd.date_range(sd, ed)
    # get price data, including SPY
    data_features = [get_data([symbol, "SPY"], dates, False)]
    # get data for the requested symbol and SPY
    columns = [symbol, "SPY"]
    # get data for each of the fields we need for the required features
    for field in self._fields:
        # columns for each field are in the form <field_name>_symbol
        data_features.append(get_data([symbol], dates, False, field)[symbol])
        columns.append(symbol + "_" + field)
        data_features.append(get_data(["SPY"], dates, False, field)["SPY"])
        columns.append("SPY_" + field)

    # combine the data into a single DataFrame
    data = pd.concat(data_features, axis=1)
    data.columns = columns
    # drop data for any date that SPY didn't trade
    data = data.dropna(subset=["SPY"])

    # create dictionary for prices (for performance reasons)
    price_dict = {}
    for row in data[symbol].iteritems():
        price_dict[row[0]] = row[1]

    # check if we've already created our indicators
    if indicators is None:
        self._indicators = Indicators(self._features)
        indicators = self._indicators

    # extract the indicators for this data set
    indicators.select(data, symbol)
    # discretize our features
    (dates, states, total_state_size) = indicators.discretize()

    # create a dictionary to store all the reachable states (for performance reasons)
    # probably not practical for a real-world system, but sufficient for this project
    state_map = {}

    # in addition to the features, there will be two other contributions to the state:
    # - two possible CR states (zero and below or greater than zero)
    # - three possible holding positions (NOT HOLDING, HOLDING LONG, HOLDING SHORT)
    # create these "sub-states" to add to our existing state data
    sub_states = []
    for i in range(2):
        for j in range(3):
            sub_states.append(str(j) + str(i))

```

```

# combine our substates with our feature data to pre-determine all reachable states
index = 0
for state in states:
    for sub_state in sub_states:
        full_state = state + sub_state
        if full_state not in state_map.keys():
            state_map[state + sub_state] = index
            index += 1

# this is the total state size (not the reachable state size)
self._print_debug("Total state size: " + str(total_state_size * 6))

return states, state_map, dates, price_dict

def _perform_trades(self, prices, symbol, sd, sv, dates, states, state_map, query_only=False):
    # set-up the portfolio
    self._portfolio = Portfolio(prices, symbol, sd, sv)

    # check if this is the first pass through
    started = False
    # number of shares we're currently holding
    share_holding = 0
    # current CR
    cr = 0.0
    # check if we should impose a penalty
    penalize = False

    for (date, state_str) in zip(dates, states):

        # get the reward for the current date (will be 0.0 to start with, and ignored on the first iteration)
        (r, cr) = self._portfolio.get_reward(date)

        # determine our current CR state
        if cr <= 0.0:
            cr_state = "0"
        else:
            cr_state = "1"

        # determine our current holding position
        if share_holding == 0:
            position = "0"
        elif share_holding > 0:
            position = "1"
        else:
            position = "2"

        # add position and cr states to get the full state-string
        state_str += (position + cr_state)

        # get the mapped state
        state = state_map[state_str]

        # we may loop the current state multiple times if the Q-Learner continues to return an invalid action
        while True:
            # if this is the first day, or we're using the learned policy only (not updating), call the appropriate
            # method
            if not started or query_only:
                # ignore states we don't know about - just hold (may occur for periods outside the training period)
                if state >= self._q_learner.num_states:
                    self._print_debug("Skipping state: " + str(state))
                    # set the action to holding
                    action = 4
            else:
                # query the Q-Learner only (don't update)
                action = self._q_learner.querysetstate(state)
                started = True

            if query_only:
                # if we've selecting an invalid action, note it (will be ignored in portfolio)
                if (action == 0) and (share_holding > 0):
                    self._print_debug("

```



Andrew 11 months ago There seems to be a limit to how much you can post in a single message. Adding the remaining code in follow-up post(s).

```

            if query_only:
                # if we've selecting an invalid action, note it (will be ignored in portfolio)
                if (action == 0) and (share_holding > 0):
                    self._print_debug("Buy 100 when holding long")
                elif (action == 1) and (share_holding < 0):
                    self._print_debug("Short 100 when holding short")
                elif (action == 2) and (share_holding >= 0):
                    self._print_debug("Buy 200 when holding none or long")
                elif (action == 3) and (share_holding <= 0):
                    self._print_debug("Short 200 when holding none short")

            else:
                # if the Q-Learner selected an invalid action, penalize it
                if penalize:

```

```

        r = -100
        # update the Q-Learner with the current state and reward, then get the next action
        action = self._q_learner.query(state, r)

    # when training the learner, penalize it if it attempts to perform an invalid action
    if not query_only:
        if (action == 0) and (share_holding > 0):
            penalize = True
        elif (action == 1) and (share_holding < 0):
            penalize = True
        elif (action == 2) and (share_holding >= 0):
            penalize = True
        elif (action == 3) and (share_holding <= 0):
            penalize = True
        else:
            penalize = False

        # if there's no penalty, we can continue
        if not penalize:
            break

    # update the portfolio and get the current share count
    share_holding = self._portfolio.update(date, self._actions[action], symbol, self._shares[action])

    # after processing all state data, return the final CR
    return cr

# simple debug method
def _print_debug(self, message):
    if self._verbose:
        print message

# copy of indicators class from previous project
class Indicators(object):

    # constructor
    def __init__(self, features):
        self._features = list(features)
        self._selected_indicators = None
        self._normalizers = None
        self._thresholds = None

    def select(self, data, symbol):

        # extract the specified features from the data set

        norms = []
        selected = []
        if not self._features or "stoch_rsi" in self._features:
            # stochastic RSI
            stoch_rsi, normalizer = Indicators._get_stochastic_rsi(data, symbol)
            selected.append(stoch_rsi)
            norms.append(normalizer)

        if not self._features or "williams_r" in self._features:
            # williams %r
            williams_r, normalizer = Indicators._get_williams_r(data, symbol)
            selected.append(williams_r)
            norms.append(normalizer)

        if not self._features or "aroon" in self._features:
            # normalized aroon oscillator (25 days)
            aroon, normalizer = Indicators._get_aroon(data, symbol)
            selected.append(aroon)
            norms.append(normalizer)

        if not self._features or "momentum" in self._features:
            # momentum (5 day)
            momentum, normalizer = \
                Indicators._get_momentum(data, symbol, 5, self._features.index("momentum"), self._normalizers)
            selected.append(momentum)
            norms.append(normalizer)

        if not self._features or "momentum_10" in self._features:
            # momentum (10 day)
            momentum, normalizer = Indicators._get_momentum(data, symbol, 10, self._features.index("momentum_10"),
                                                             self._normalizers)
            selected.append(momentum)
            norms.append(normalizer)

        if not self._features or "momentum_20" in self._features:
            # momentum (20 day)
            momentum, normalizer = Indicators._get_momentum(data, symbol, 20, self._features.index("momentum_20"),
                                                             self._normalizers)
            selected.append(momentum)
            norms.append(normalizer)

        if not self._features or "volume" in self._features:
            # volume
            volume, normalizer = Indicators._get_volume(data, symbol, self._features.index("volume"), self._normalizers)

```

```

        selected.append(volume)
        norms.append(normalizer)

    if not self._features or "volatility" in self._features:
        # volatility
        volatility, normalizer = \
            Indicators._get_volatility(data, symbol, self._features.index("volatility"), self._normalizers)
        selected.append(volatility)
        norms.append(normalizer)

    if not self._features or "5_day" in self._features:
        # 5_day ewma
        five_day, normalizer = Indicators._get_normalized_ema(data, symbol, 5)
        selected.append(five_day)
        norms.append(normalizer)

    if not self._features or "10_day" in self._features:
        # 10_day ewma
        ten_day, normalizer = Indicators._get_normalized_ema(data, symbol, 10)
        selected.append(ten_day)
        norms.append(normalizer)

    if not self._features or "20_day" in self._features:
        # 20_day ewma
        twenty_day, normalizer = Indicators._get_normalized_ema(data, symbol, 20)
        selected.append(twenty_day)
        norms.append(normalizer)

    if not self._features or "5_day_spy" in self._features:
        # 5_day_spy ewma
        twenty_day_spy, normalizer = Indicators._get_normalized_ema(data, "SPY", 5)
        selected.append(twenty_day_spy)
        norms.append(normalizer)

    if not self._features or "10_day_spy" in self._features:
        # 10_day_spy ewma
        twenty_day_spy, normalizer = Indicators._get_normalized_ema(data, "SPY", 10)
        selected.append(twenty_day_spy)
        norms.append(normalizer)

    if not self._features or "20_day_spy" in self._features:
        # 20_day_spy ewma
        twenty_day_spy, normalizer = Indicators._get_normalized_ema(data, "SPY", 20)
        selected.append(twenty_day_spy)
        norms.append(normalizer)

    indicators = pd.concat(selected, axis=1).dropna()
    indicators.columns = self._features

    self._selected_indicators = indicators
    self._normalizers = norms

@property
def selected_indicators(self):
    return self._selected_indicators

@property
def normalizers(self):
    return self._normalizers

@property
def features(self):
    return self._features

def discretize(self):
    # to maintain consistency for all runs, make sure we use the same thresholds
    if self._thresholds is None:
        # no thresholds - find them
        self.find_thresholds()

    binned = []
    total_state_size = 1
    for feature in self._features:
        # get thresholds for the current feature
        thresholds = self._thresholds.get(feature)
        # we have one less bin than number of thresholds
        num_bins = len(thresholds) - 1
        # hack to deal with lack of volume data for ML4T-220
        if num_bins == 0:
            num_bins = 1

        # increment the total state size
        total_state_size *= num_bins

        # ensure all features are within the outer bounds
        self._selected_indicators[feature][self._selected_indicators[feature] < thresholds[0]] = thresholds[0]
        self._selected_indicators[feature][self._selected_indicators[feature] > thresholds[len(thresholds) - 1]] = \
            thresholds[len(thresholds) - 1]

        # just label our bins with the bin number
        labels = [i for i in range(num_bins)]

```

```

        # discretize the features
        (bin_features, feature_thresholds) = pd.cut(self._selected_indicators[feature], bins=thresholds,
                                                    labels=labels, right=True, include_lowest=True, retbins=True)

        binned.append(bin_features)

    # combine discretized features
    binned_ind = pd.concat(binned, axis=1)

    # combine the discretized features into a single state string for each day
    states = []
    for row in binned_ind.itertuples():
        state_str = ""
        for i in range(1, len(self._features) + 1):
            state_str += str(row[i])
        states.append(state_str)

    # return the dates, states and the total state size
    return binned_ind.index, states, total_state_size

def find_thresholds(self):
    # set-up the threshold dictionary for each of the features
    self._thresholds = {}
    for feature in self._features:
        # find the number of unique values for the given feature
        num_unique = len(self._selected_indicators[feature].unique())
        # just use simple log of the number of unique values to determine the number of bins to use
        num_bins = int(math.log(num_unique, 2))
        # hack to deal with missing volume field for ML4T-220
        if num_bins == 0:
            num_bins = 1

        # just label our bins with the bin number
        labels = [i for i in range(num_bins)]

        # discretize the features using the number of bins (equal-width bins)
        (bin_features, feature_thresholds) = pd.cut(self._selected_indicators[feature]

```



Andrew 11 months ago And again...

```

        # discretize the features using the number of bins (equal-width bins)
        (bin_features, feature_thresholds) = pd.cut(self._selected_indicators[feature], bins=num_bins,
                                                    labels=labels, right=True, include_lowest=True, retbins=True)

        # set the thresholds for this feature
        self._thresholds[feature] = feature_thresholds

    @staticmethod
    def _get_stochastic_rsi(data, symbol):
        # calculate RSI from gains/losses
        gain_or_loss = data.diff()[symbol]
        gains, losses = gain_or_loss.copy(), gain_or_loss.copy()

        # calculate average gains over the window
        gains[gain_or_loss < 0] = 0
        avg_gain = pd.rolling_mean(gains, window=14)

        # calculate average losses over the window
        losses[gain_or_loss > 0] = 0
        # use absolute value of losses
        losses = losses.abs()
        avg_loss = pd.rolling_mean(losses, window=14)

        rs = (avg_gain / avg_loss)
        rsi = 100.0 - (100.0 / (1.0 + rs))

        return pd.rolling_apply(rsi, window=14, func=Indicators._calculate_stoch_rsi), None

    @staticmethod
    def _calculate_stoch_rsi(x):
        return (x[len(x) - 1] - x.min()) / (x.max() - x.min()) - 0.5

    @staticmethod
    def _get_momentum(data, symbol, window, index, normalizers):
        momentum = data[symbol]/data[symbol].shift(window)
        return Indicators._normalize_data(momentum, index, normalizers)

    @staticmethod
    def _get_aroon(data, symbol):
        aroon_up = pd.rolling_apply(data[symbol], window=25, func=Indicators._calculate_aroon_up)
        aroon_down = pd.rolling_apply(data[symbol], window=25, func=Indicators._calculate_aroon_down)

        return aroon_up - aroon_down - 0.5, None

    @staticmethod
    def _calculate_aroon_up(x):
        return (25 - np.argmax(x))/25.0

```



```

@staticmethod
def _calculate_arrron_down(x):
    return (25 - np.argmax(x))/25.0

@staticmethod
def _get_williams_r(data, symbol):
    index = np.array(range(len(data.index)))
    williams_r = pd.rolling_apply(index, window=14,
                                  func=lambda x: Indicators._calculate_williams_r(x, data, symbol))
    cur = pd.DataFrame(data=williams_r, index=data.index)
    return cur, None

@staticmethod
def _calculate_williams_r(x, data, symbol):
    cur_loc = int(x[len(x) - 1])
    cur_close = data.iloc[cur_loc][symbol + "_Close"]

    highest_high = data.iloc[x][symbol + "_High"].max()
    lowest_low = data.iloc[x][symbol + "_Low"].min()
    return (highest_high - cur_close) / (highest_high - lowest_low) - 0.5

@staticmethod
def _get_volume(data, symbol, index, normalizers):
    if data[symbol + "_Volume"].ix[1] == 1:
        return data[symbol + "_Volume"], None
    return Indicators._normalize_data(data[symbol + "_Volume"], index, normalizers)

@staticmethod
def _normalize_data(data, index, normalizers):
    normalizer = Normalizer(data.mean(), data.std())
    if normalizers is not None:
        normalizer = normalizers[index]

    return normalizer.normalize(data), normalizer

@staticmethod
def _get_volatility(data, symbol, index, normalizers):
    return Indicators._normalize_data(pd.rolling_std(data[symbol], window=10), index, normalizers)

@staticmethod
def _get_normalized_ema(data, symbol, window):
    sma = pd.ewma(data[symbol], span=window)
    std = pd.rolling_std(data[symbol], window=window)
    return (data[symbol] - sma) / (2 * std), None

class Normalizer(object):

    # constructor
    def __init__(self, mean, std):
        self._mean = mean
        self._std = std

    def normalize(self, values):
        return (values - self._mean) / (2 * self._std)

class Portfolio(object):

    def __init__(self, prices, symbol, start_date, sv=10000):
        self._prices = prices
        self._start_date = start_date
        self._sv = sv
        self._symbols = [symbol]

        # min/max shares allowed for any given day
        self._max_shares = 100
        self._min_shares = -100

        # starting positions
        self._current_positions = {"CASH": self._sv, symbol: 0}

        # track trades in a dictionary
        self._trades = {}

        # track the last CR calculated
        self._last_cr = sv

    def update(self, date, action, symbol, shares):

        # if the action is hold, do nothing

        # if the action is buy, update the portfolio
        if action == "BUY":
            # make sure it's a legal action
            if (self._current_positions[symbol] + shares) <= self._max_shares:
                # update number of shares
                self._current_positions[symbol] += shares
                # update current cash
                self._current_positions["CASH"] = self._current_positions["CASH"] - shares * self._prices[date]

```

```

        # update trades dictionary
        self._trades[date] = shares

# if the action is sell, update the portfolio
elif action == "SELL":
    # make sure it's a legal action
    if (self._current_positions[symbol] - shares) >= self._min_shares:
        # update number of shares
        self._current_positions[symbol] -= shares
        # update current cash
        self._current_positions["CASH"] = self._current_positions["CASH"] + shares * self._prices[date]
        # update trades dictionary
        self._trades[date] = -shares

# return the current asset position
return self._current_positions[symbol]

def get_reward(self, date):
    # get the reward for this date based on daily returns
    # first, simply calculate the current portfolio value
    current_totals = 0
    for entry in self._current_positions.keys():
        if entry == "CASH":
            current_totals += self._current_positions[entry]
        else:
            current_totals += self._current_positions[entry] * self._prices[date]

    # daily returns is based on the relationship of the current portfolio value and yesterdays portfolio value
    daily_rets = (current_totals / self._last_cr) - 1

    # cr is based on the current portfolio value and the starting value
    cr = (current_totals / self._sv) - 1

    # update the Last CR
    self._last_cr = current_totals

    # return the reward and the CR
    return daily_rets, cr

@property
def trades(self):
    # generate a trades DataFrame from the stored trades dictionary
    index = np.array(sorted(self._prices.keys()))

    trades_df = pd.DataFrame(index=index, columns=self._symbols)
    trades_df = trades_df.fillna(value=0.)
    for key in self._trades.keys():
        # we're only training/testing on one symbol
        trades_df[self._symbols[0]][key] = self._trades[key]

    return trades_df

```



Andrew 11 months ago I tried to start the each post slightly overlapping the previous...hopefully it's not too hard to follow.